

Natural Language Processing (NLP) and Large Language Models (LLMs)

Lecture 3-1: Machine learning II

Chendi Wang (王晨笛)
chendi.wang@xmu.edu.cn

WISE @ XMU

2025 年 3 月 13 日

- ① Section 1: Classification and NLP
- ② Section 2: Hyperparameter tuning
- ③ Section 3: Adam
- ④ Section 4: Normalization
- ⑤ Section 5: Convolutional neural networks

① Section 1: Classification and NLP

② Section 2: Hyperparameter tuning

③ Section 3: Adam

④ Section 4: Normalization

⑤ Section 5: Convolutional neural networks

Many NLP Tasks are Classification Problems

- Classify sentences into positive, negative, or neutral sentiment (sentiment analysis).
- Given a center word and its surrounding context, classify the center word into categories such as Person, Location, Organization, or Other.

Named-Entity Recognition

- **Named-Entity Recognition (NER)**: find and classify names in text.
- **Kobe Bryant** was a legendary basketball player for the **Los Angeles Lakers** from **1996 to 2016**, after completing his high school education in **Philadelphia**.
 - **Kobe Bryant**: Person
 - **Los Angeles Lakers**: Organization
 - **1996 to 2016**: Time
 - **Philadelphia**: Location

Classification notation

- Input: $\mathbf{x} \in \mathbb{R}^d$ is a word vector, or a sentence, or a document etc.
- Label: $y \in \{1, 2, \dots, K\}$
 - sentiment: (1= positive, 2= negative)
 - named entities, 1 = Person; 2= Time, ...
 - decision: 1=buy, 2= sell
- Onehot label: $\mathbf{y} = [y^{(1)}, \dots, y^{(K)}] = [0, \dots, 0, 1, 0, \dots, 0]$ ($y^{(k)} = 1$ if $y = k$).

Window Classification

- Goal: **Classify each word based on its context window.**
- Consider a sequence of word embeddings within a sentence:

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(T)}\}, \quad \text{each } \mathbf{x}^{(t)} \in \mathbb{R}^d.$$

- For each word $\mathbf{x}^{(t)}$, define a context window of size $2K + 1$:

$$\mathbf{x}^{(t-K)}, \dots, \mathbf{x}^{(t-1)}, \mathbf{x}^{(t)}, \mathbf{x}^{(t+1)}, \dots, \mathbf{x}^{(t+K)}.$$

- Construct input vector by concatenating embeddings in the window:

$$\mathbf{x}_t = [\mathbf{x}^{(t-K)}, \dots, \mathbf{x}^{(t-1)}, \mathbf{x}^{(t)}, \mathbf{x}^{(t+1)}, \dots, \mathbf{x}^{(t+K)}] \in \mathbb{R}^{(2K+1)d}.$$

- Label: y_t is the class label for the **center word** $\mathbf{x}^{(t)}$.
- Training data: $\{(\mathbf{x}_t, y_t)\}_{t=1}^T$

Training

- Training data: $\{\mathbf{z}_i = (\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^n$.
- For a classifier (such as a neural network or a linear classifier) $\mathbf{f}_\theta = (f_\theta^{(1)}, \dots, f_\theta^{(K)})$, **find θ that minimizes:**

$$\frac{1}{n} \sum_{i=1}^n \ell_\theta^{\text{CE}}(\mathbf{z}_i) = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^K y_i^{(j)} \log f_\theta^{(j)}(\mathbf{x}_i).$$

Neural networks update the embeddings

- Recall

$$\mathbf{f}_{\theta}(\mathbf{x}) = \text{SoftMax} \left(\mathbf{W}^{[L]} \sigma(\mathbf{W}^{[L-1]} \sigma(\dots \sigma(\mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]})) + \mathbf{b}^{[L-1]}) \right).$$

- Rewrite

$$\mathbf{f}_{\theta}(\mathbf{x}) = \text{SoftMax} \left(\mathbf{W}^{[L]} \mathbf{h}(\mathbf{x}) \right)$$

with $\mathbf{h}(\mathbf{x}) = \sigma(\mathbf{W}^{[L-1]} \sigma(\dots \sigma(\mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]})) + \mathbf{b}^{[L-1]})$.

- $\mathbf{h}(\mathbf{x}) \in \mathbb{R}^{d^{[L]}}$ is a new word vector (a.k.a., embedding, representation, or feature for image data).
- Note that \mathbf{h} is updated when θ is updated.


Training

- Training data: $\{\mathbf{z}_i = (\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^n$.
- For a classifier (such as a neural network or a linear classifier) $\mathbf{f}_\theta = (f_\theta^{(1)}, \dots, f_\theta^{(K)})$, **find θ that minimizes:**


$$\frac{1}{n} \sum_{i=1}^n \ell_\theta^{\text{CE}}(\mathbf{z}_i) = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^K y_i^{(j)} \log f_\theta^{(j)}(\mathbf{x}_i).$$

Retraining word embeddings

- The pretrained word vectors (such as using GloVe) could be further **retrained in downstream tasks** (such as classification) to perform better.
- However, retraining is risky (see the figure).
- When retraining, we need to ensure that the **training set is large enough to cover most words from the vocabulary**.



Linear classifiers before retraining



Linear classifier after retraining

An example when linear classifiers fail to classify the word “television” after retraining

- ① Section 1: Classification and NLP
- ② Section 2: Hyperparameter tuning
- ③ Section 3: Adam
- ④ Section 4: Normalization
- ⑤ Section 5: Convolutional neural networks

Recall: Ridge regression

- Linear ridge regression:

$$\min_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n \left(y_i - \mathbf{w}^T \mathbf{x}_i - b \right) + \lambda \sum_{j=1}^d w_j^2$$

- $\lambda \sum_{j=1}^d w_j^2$ is a regularization term that is used to prevent overfitting.
- λ here is a **hyperparameter** (adjusting bias-variance trade-off).

Overfitting and underfitting



Figure is from <https://www.geeksforgeeks.org/underfitting-and-overfitting-in-machine-learning/>

Bias-variance trade-off

- For a data point (\mathbf{x}, y) such that $y = f_*(\mathbf{x}) + \epsilon$ with $\mathbb{E}[\epsilon|\mathbf{x}] = 0$, consider an estimate $\hat{y} = f_{\hat{\theta}}(\mathbf{x})$ ($\hat{y} = \hat{\mathbf{w}}^T \mathbf{x} + \hat{b}$ for linear models).
- $\hat{\theta}$ is random since it is trained on training data.
- $\mathbb{E}_{\hat{\theta}, \epsilon}[(y - f_{\hat{\theta}}(\mathbf{x}))^2] = \text{Bias}^2 + \text{Variance} + \sigma^2$.
- $\sigma^2 = \mathbb{E}[\epsilon^2|\mathbf{x}]$ is called the Bayes error (lowest error rate).
- Bias: $\mathbb{E}_{\hat{\theta}}[\hat{y}] - f_*(\mathbf{x})$
- Variance: $\text{Var}[\hat{y}] = \mathbb{E}_{\hat{\theta}}[\hat{y} - \mathbb{E}_{\hat{\theta}}[\hat{y}]]^2$.
- Bias increases, variance decreases (and vice versa); controlled by λ .
- One can verify that OLS is un-biased, check the variance (and calculate the bias and variance for linear ridge regression).

Regularization in deep learning

- In deep learning, we usually don't use **explicit regularization**.
- **Implicit regularization** can be achieved by **tuning hyperparameters**.
- Trainable parameters (estimated by training data): θ
- Hyperparameters can **not** be estimated using **training data**.

An example of training a convolutional neural network (CNN)

- train-CNN.ipynb
- device:CPU v.s. GPU
- An epoch is a complete pass through the **entire training dataset** during the training of a machine learning model.

Running time for CPU

```
Epoch 1, Iteration 1: Time = 0.3002 seconds
Epoch 1, Iteration 2: Time = 0.0606 seconds
Epoch 1, Iteration 3: Time = 0.0610 seconds
Epoch 1, Iteration 4: Time = 0.0704 seconds
Epoch 1, Iteration 5: Time = 0.0659 seconds
Epoch 1, Iteration 6: Time = 0.0602 seconds
Epoch 1, Iteration 7: Time = 0.0578 seconds
Epoch 1, Iteration 8: Time = 0.0603 seconds
Epoch 1, Iteration 9: Time = 0.0684 seconds
Epoch 1, Iteration 10: Time = 0.0651 seconds
Epoch 1, Iteration 11: Time = 0.0592 seconds
Epoch 1, Iteration 12: Time = 0.0651 seconds
Epoch 1, Iteration 13: Time = 0.0656 seconds
Epoch 1, Iteration 14: Time = 0.0676 seconds
Epoch 1, Iteration 15: Time = 0.0678 seconds
Epoch 1, Iteration 16: Time = 0.0657 seconds
Epoch 1, Iteration 17: Time = 0.0658 seconds
Epoch 1, Iteration 18: Time = 0.0664 seconds
```

Running time for GPU

```
Epoch 1, Iteration 1: Time = 0.9922 seconds
Epoch 1, Iteration 2: Time = 0.0057 seconds
Epoch 1, Iteration 3: Time = 0.0036 seconds
Epoch 1, Iteration 4: Time = 0.0040 seconds
Epoch 1, Iteration 5: Time = 0.0035 seconds
Epoch 1, Iteration 6: Time = 0.0036 seconds
Epoch 1, Iteration 7: Time = 0.0035 seconds
Epoch 1, Iteration 8: Time = 0.0035 seconds
Epoch 1, Iteration 9: Time = 0.0035 seconds
Epoch 1, Iteration 10: Time = 0.0035 seconds
Epoch 1, Iteration 11: Time = 0.0045 seconds
Epoch 1, Iteration 12: Time = 0.0036 seconds
Epoch 1, Iteration 13: Time = 0.0035 seconds
Epoch 1, Iteration 14: Time = 0.0035 seconds
Epoch 1, Iteration 15: Time = 0.0036 seconds
Epoch 1, Iteration 16: Time = 0.0035 seconds
Epoch 1, Iteration 17: Time = 0.0036 seconds
Epoch 1, Iteration 18: Time = 0.0035 seconds
```

Some hyperparameters

- Learning rate η
- Mini-batch size b (related to learning rate)
- Depth L
- Number of neurons per layer $d^{[l]}, l = 1, \dots, L$.
- Algorithms: SGD, Adam, Nesterov acceleration etc.
- Number of iterations (early stopping)
- Dropout (set some parameters to be 0)
- Initialization
- Temperature (for generative models) etc.

Initialization

- One simple way is to initialize parameters as $\theta_0 = 0$.
- In train-CNN.ipynb, we use the default method: **random initialization**, also known as **Kaiming (He) initialization**.
- **Random Initialization**: For each layer's weight matrix

$$\mathbf{W}^{[l]} = [\mathbf{w}_1^{[l]}, \dots, \mathbf{w}_{d^{[l+1]}}^{[l]}] \in \mathbb{R}^{d^{[l]} \times d^{[l+1]}},$$

initialize each column vector independently as

$$\mathbf{w}_i^{[l]} \sim \mathcal{N}(0, \sigma^2 I), \quad \text{with} \quad \sigma^2 = \frac{2}{d^{[l]}}.$$

- Bias terms $\mathbf{b}^{[l]}$ are typically initialized as $\mathbf{0}$.

Early Stopping in Deep Learning

- **Idea:** Stop training when performance on a validation set no longer improves, to prevent **overfitting**.
- **Procedure:** Track validation loss during training; halt training when validation loss starts increasing or remains unchanged. (Example early-stopping.ipynb)
- **Benefit:** Provides an automatic and practical way to select the number of epochs, improving generalization.
- One can show that early stopping is mathematically similar to ridge regression for linear models.

Dropout

- Consider the hidden layer output $\mathbf{h} \in \mathbb{R}^{d^{[l]}}$ (e.g., $\mathbf{h} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$ for $l = 1$).
- Define the random masking vector $\mathbf{r} = [r_1, \dots, r_{d^{[l]}}] \in \{0, 1\}^{d^{[l]}}$, where each r_i independently satisfies:

$$\text{Prob}[r_i = 0] = p_{\text{drop}}.$$

- Dropout applies as:

$$\mathbf{h}_{\text{drop}} = \gamma(\mathbf{r} \odot \mathbf{h}),$$

where \odot denotes the Hadamard (element-wise) product.

- The scaling factor γ ensures unbiasedness:

$$\gamma \mathbb{E}[r_i h_i] = h_i \quad \text{for all } i.$$

- Dropout helps prevent overfitting by randomly setting a fraction of activations to zero, reducing co-adaptation among neurons.

Dropout in pytorch

```
▶ import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F

▶ class SimpleNet(nn.Module):
    def __init__(self):
        super(SimpleNet, self).__init__()
        self.fc1 = nn.Linear(784, 256) # Input layer to hidden layer
        self.dropout1 = nn.Dropout(p=0.5) # Dropout layer with 50% probability
        self.fc2 = nn.Linear(256, 128) # Hidden layer to hidden layer
        self.dropout2 = nn.Dropout(0.5) # Dropout layer with 50% probability
        self.fc3 = nn.Linear(128, 10) # Hidden layer to output layer

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.dropout1(x) # Apply dropout
        x = F.relu(self.fc2(x))
        x = self.dropout2(x)
        x = self.fc3(x)
        return x
```


Learning rate

- The magnitude of the learning rate matters!!!!
- Large learning rate: may not converge
- Small learning rate: converges slowly
- No fixed rule, depending on your task and model (Experimentation and observation are key)
- A common combination: warmup + decay + adaptive methods.

Why is Warm-up Important?

- In the early stages of training, gradients are typically large.
- If the learning rate is high initially, parameters may change dramatically, causing training instability.
- Gradually increasing the learning rate (warm-up) helps stabilize training and enables better convergence.
- It also assists in escaping poor local minima.

What is Warm-up?

- We start from a small learning rate η_{init} and increase it to a relatively large η_{base} in the **first T_{warm} steps (or epochs)**.
- Mathematically, for the t -th step we have

$$\eta_t = \eta_{\text{init}} + (\eta_{\text{base}} - \eta_{\text{init}}) \cdot \frac{t}{T_{\text{warm}}}, \quad 0 \leq t \leq T_{\text{warm}}.$$

- See `warmup.ipynb`

Learning rate decay (after warm-up)

```
trainer = torch.optim.SGD(net.parameters(), lr=0.1)
scheduler = lr_scheduler.StepLR(trainer, step_size=5, gamma=0.1)
```

- $lr = \eta_{\text{base}}$
- Polynomial decay: the learning rate of Epoch t is $lr \times \text{gamma}^{\lfloor t/\text{step_size} \rfloor}$
- Other decay rules: exponential decay, Cosine Annealing etc.

Optimized Hyperparameter Tuning: Grid Search

- Grid Search: Evaluates all combinations from a specified set of parameters.
- Advantages: simple to implement and interpret (cause it is optimized over the defined grid); easy parallelization.
- Disadvantages: Computationally expensive (time-consuming); Efficiency decreases exponentially with the number of parameters.
- Common libraries: `scikit-learn`, `Optuna`.


- ① Section 1: Classification and NLP
- ② Section 2: Hyperparameter tuning
- ③ Section 3: Adam
- ④ Section 4: Normalization
- ⑤ Section 5: Convolutional neural networks

GD may not converge

- Minimize

$$g(\theta_1, \theta_2) = 0.1\theta_1^2 + 2\theta_2^2.$$

- See gd-counter-example.ipynb
- The gradient in the θ_2 direction is much larger and is changing more rapidly than the horizontal direction (θ_1).



Example is from https://d2l.ai/chapter_optimization/momentum.html

Recap: SGD

- Loss function:

$$\frac{1}{n} \sum_{i=1}^n \ell_{\theta}(\mathbf{z}_i).$$

- In the t -th iteration, sample $\mathcal{B}_t \subset \{1, 2, \dots, n\}$ with mini-batch size $|\mathcal{B}_t| = b$
- SGD update $\theta_{t+1} = \theta_t - \frac{\eta_t}{b} \sum_{i \in \mathcal{B}_t} \nabla_{\theta} \ell_{\theta}(\mathbf{z}_i)$.
- Rewrite $\mathbf{g}_t = \frac{1}{b} \sum_{i \in \mathcal{B}_t} \nabla_{\theta} \ell_{\theta}(\mathbf{z}_i)$ and

$$\theta_{t+1} = \theta_t - \eta_t \mathbf{g}_t.$$

The momentum method

- Let

$$\mathbf{v}_{t+1} = \beta \mathbf{v}_t + \mathbf{g}_t.$$

- Updates

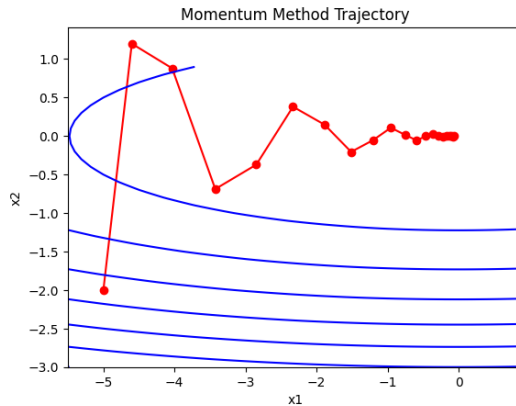
$$\theta_{t+1} = \theta_t - \eta \mathbf{v}_{t+1}$$

- Momentum replaces gradients with a leaky average over past gradients, that is

$$\mathbf{v}_{t+1} = \sum_{s=0}^{t-1} \beta^s \mathbf{g}_s.$$

Example

- Recall the example in `gd-counter-example.ipynb`.
- Using momentum method, we have



Adagrad (Duchi et al., 2011)

- Adagrad adapts the learning rate by accumulating **past squared gradients** using a state vector \mathbf{s}_t .
- Denote the gradient at step t by $\mathbf{g}_t = \frac{1}{b} \sum_{j \in \mathcal{B}_t} \nabla_{\theta} \ell_{\theta}(\mathbf{z}_j)$.

- **Iteration formula:**

$$\mathbf{s}_t = \mathbf{s}_{t-1} + \mathbf{g}_t^2, \quad \theta_t = \theta_{t-1} - \frac{\eta \mathbf{g}_t}{\sqrt{\mathbf{s}_t + \delta \cdot \mathbf{1}}}$$

(All operations are element-wise.)

- Here, δ is a small constant for numerical stability (typically 10^{-6} or 10^{-7}).
- Example `gd-counter-example.ipynb` (adjusting learning rate)

RMSProp (Tieleman and Hinton, 2012)

- In the original Adagrad algorithm, the accumulator vector \mathbf{s}_t grows without bound, causing the learning rate to decay excessively.
- A common modification is to introduce a decay factor $0 < \gamma < 1$:

$$\mathbf{s}_t = \gamma \mathbf{s}_{t-1} + (1 - \gamma) \mathbf{g}_t^2, \quad \theta_t = \theta_{t-1} - \frac{\eta \mathbf{g}_t}{\sqrt{\mathbf{s}_t + \delta \cdot \mathbf{1}}}.$$

(Element-wise operations)

- Example implementation: `gd-counter-example.ipynb`

Adam (Kingma and Ba; 2015)

- Adam combines the ideas of momentum and RMSProp:

$$\text{Momentum: } \mathbf{v}_t = \beta \mathbf{v}_{t-1} + (1 - \beta) \mathbf{g}_t$$

$$\text{RMSProp: } \mathbf{s}_t = \gamma \mathbf{s}_{t-1} + (1 - \gamma) \mathbf{g}_t^2$$

- Bias correction step (to offset initialization at zero):

$$\hat{\mathbf{v}}_t = \frac{\mathbf{v}_t}{1 - \beta^t}, \quad \hat{\mathbf{s}}_t = \frac{\mathbf{s}_t}{1 - \gamma^t}$$

(since $1 + \beta + \beta^2 + \dots + \beta^{t-1} = \frac{1 - \beta^t}{1 - \beta}$)

- Parameter update rule:

$$\theta_t = \theta_{t-1} - \frac{\eta_t \hat{\mathbf{v}}_t}{\sqrt{\hat{\mathbf{s}}_t} + \delta \cdot \mathbf{1}}$$

(all operations element-wise)

- Example implementation: `gd-counter-example.ipynb`

Adam (Kingma and Ba; 2015)

Train a CNN using Adam: Adam-CNN.ipynb

Weight decay (l_2 -penalty)

- l_2 -penalty: $\frac{1}{n} \sum_{i=1}^n \ell_{\theta}(\mathbf{z}_i) + \lambda \|\theta\|_2^2$
- SGD with l_2 -penalty

$$\theta_t = \theta_{t-1} - \eta (\mathbf{g}_t + 2\lambda\theta_{t-1}).$$

- This is also known as weight decay as the norm of θ decays rapidly due to the l_2 -penalty.

AdamW: Adam with Decoupled Weight Decay Regularization

- Standard l_2 regularization is not equivalent to weight decay in Adam.
- AdamW explicitly decouples weight decay from the gradient update:

$$\theta_t = \theta_{t-1} - 2\lambda\theta_{t-1} - \frac{\eta_t \hat{\mathbf{v}}_t}{\sqrt{\hat{\mathbf{s}}_t} + \delta \cdot \mathbf{1}}$$

- AdamW typically achieves better performance across various language modeling tasks.

```
optimizer = optim.AdamW(model.parameters(), lr=0.001, weight_decay=0.01)
```


- ① Section 1: Classification and NLP
- ② Section 2: Hyperparameter tuning
- ③ Section 3: Adam
- ④ Section 4: Normalization
- ⑤ Section 5: Convolutional neural networks

Batch normalization (Ioffe and Szegedy, 2015)

- Motivation: solve a covariate shift problem (the input features (covariates) of the training data and testing data are different).

Batch normalization

- For each input feature \mathbf{x} , let $\mathbf{h}^{l-1}(\mathbf{x}) = \sigma(\mathbf{W}^{l-1}\sigma(\dots\sigma(\mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^1)) + \mathbf{b}^{l-1}) \in \mathbb{R}^{d^{l-1}}$ be the output of the $(l-1)$ -th layer.
- For a minibatch \mathcal{B}_t of size b , denote $\mathbf{a}_i = \mathbf{W}^l \mathbf{h}^{l-1}(\mathbf{x}_i) + \mathbf{b}^l$ for any $i \in \mathcal{B}_t$.
- Mean: $\boldsymbol{\mu} = \frac{1}{b} \sum_{i \in \mathcal{B}_t} \mathbf{a}_i$; Variance: $\boldsymbol{\sigma}^2 = \frac{1}{b} \sum_{i \in \mathcal{B}_t} (\mathbf{a}_i - \boldsymbol{\mu}) \odot (\mathbf{a}_i - \boldsymbol{\mu})$
- Normalization:

$$\tilde{\mathbf{a}}_i = \frac{\mathbf{a}_i - \boldsymbol{\mu}}{\sqrt{\boldsymbol{\sigma}^2 + \delta}},$$

$\delta > 0$ is a small number to prevent 0 values.

Batch normalization

- Introduce two more (trainable) parameter vectors γ, β to allow for heterogeneity.
- For each $i \in \mathcal{B}_t$, output: $\gamma \odot \tilde{\mathbf{a}}_i + \beta$
- Activation:

$$\sigma(\gamma \odot \tilde{\mathbf{a}}_i + \beta)$$

Batch normalization using pytorch

```
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(784, 256)
        self.bn1 = nn.BatchNorm1d(256)
        self.fc2 = nn.Linear(256, 128)
        self.bn2 = nn.BatchNorm1d(128)
        self.fc3 = nn.Linear(128, 10)

    def forward(self, x):
        x = F.relu(self.bn1(self.fc1(x)))
        x = F.relu(self.bn2(self.fc2(x)))
        x = self.fc3(x)
        return x
```

Layer normalization

- Batch normalization (BN) depends on the mini-batch size.
- It is not obvious how to use BN in recurrent neural networks (RNN).
- RNN is important to model language tasks (next week).

Layer normalization

- Batch normalization (BN) depends on the mini-batch size.
- It is not obvious how to use BN in recurrent neural networks (RNN).
- RNN is important to model language tasks (next week).

Layer normalization

- $\mathbf{a}_i = \mathbf{W}^l \mathbf{h}^{l-1}(\mathbf{x}_i) + \mathbf{b}^l \in \mathbb{R}^{d^l} =: (a_{i1}, \dots, a_{id^l})$.
- Mean: $\mu = \frac{1}{d^l} \sum_{j=1}^{d^l} a_{ij}$; Variance: $\sigma^2 = \frac{1}{d^l} \sum_{j=1}^{d^l} (a_{ij} - \mu)^2$

Layer normalization using pytorch

```
class SimpleNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.ln1 = nn.LayerNorm(hidden_size) # Layer Normalization after first layer
        self.fc2 = nn.Linear(hidden_size, output_size)
        self.ln2 = nn.LayerNorm(output_size) # Layer Normalization after second layer
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.ln1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.ln2(x)
        return x
```

- ① Section 1: Classification and NLP
- ② Section 2: Hyperparameter tuning
- ③ Section 3: Adam
- ④ Section 4: Normalization
- ⑤ Section 5: Convolutional neural networks

Code

```
nn.Conv2d(1, 6, kernel_size=5, padding=2)
```

- Input channel size: 1
- Output channel size:6
- Kernel size:5
- Padding:2

Refers to Dive into Deep Learning