

# Natural Language Processing (NLP) and Large Language Models (LLMs)

## Lecture 8-1: Pretraining

Chendi Wang (王晨笛)  
chendi.wang@xmu.edu.cn

WISE @ XMU

2025 年 4 月 22 日

① Section 1: An overview of pretraining and finetuning

② Section 2: Details of pretraining

③ Section 3: Details of GPT and BERT



# The age of pretraining

- BERT: Devlin et al., 2018
- GPT-2: Radford et al., 2019
- GPT-3: Brown et al., 2020
- Vision transformer
- GPT-4, deepseek, Llama, qwen, ...

# Pretraining

- Pretraining generally refers to training a very large model on a very large dataset.
- Empirical scaling laws demonstrate the relationship between model size, dataset size, and performance: Kaplan et al., 2020




Figure is from Kaplan et al., 2020

# Pretrained Models Achieve Strong Performance

## Benchmarks

Add a Result

These leaderboards are used to track progress in Language Modelling

Trend	Dataset	Best Model	Paper	Code	Compare
	WikiText-103	RETRO (7.5B)			See all
	Penn Treebank (Word Level)	GPT-3 (Zero-Shot)			See all
	enwik8	GPT-2 (48 layers, h=1600)			See all
	The Pile	Test-Time Fine-Tuning with SIFT + Llama-3.2 (3B)			See all
	WikiText-2	SparseGPT (175B, 50% Sparsity)			See all
	LAMBADA	PaLM-540B (Few-Shot)			See all
	One Billion Word	MDLM (AR baseline)			See all
	Text8	GPT-2			See all

## Foundation Models

- Models such as GPT-3, GPT-2, and BERT are referred to as **foundation models** due to their critically central yet inherently incomplete nature (Bommasani et al., 2021).
- A foundation model is any model **pretrained on broad data**—typically using large-scale self-supervision—that can be **adapted (e.g., via fine-tuning) to a wide range of downstream tasks**.

# Pretrained foundation models

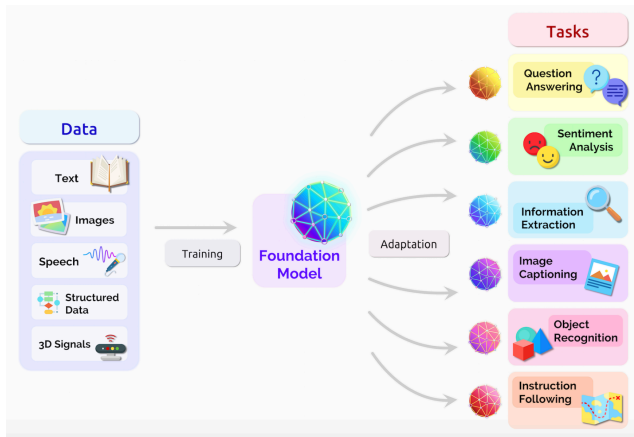



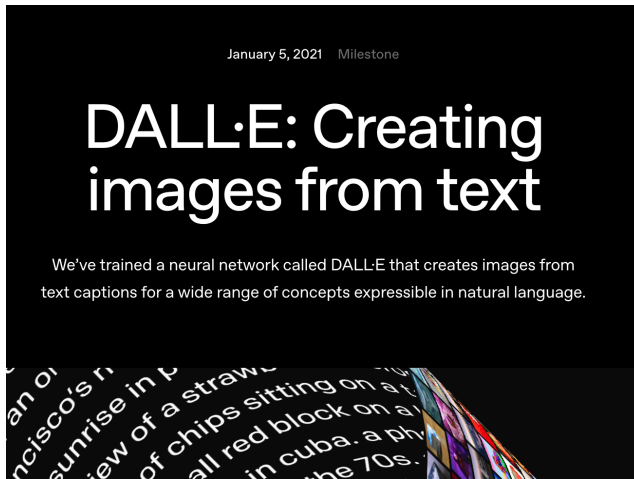
Figure is from Bommasani et al., 2021



# Foundation models are not new to us.



## Foundation models can be multimodal



## What can pretrained LLMs do?

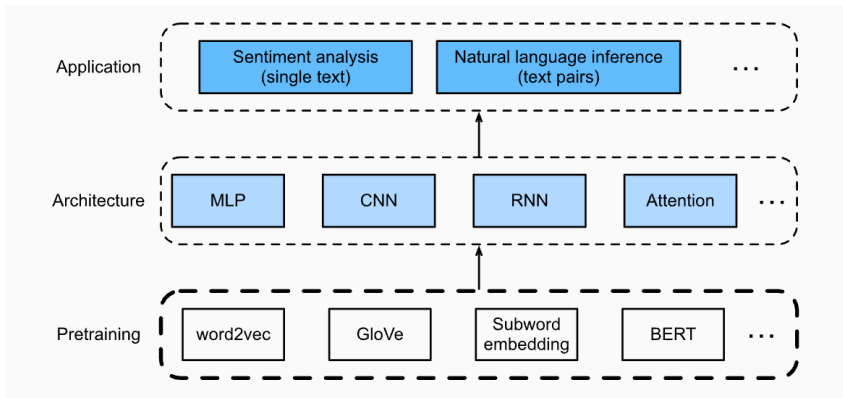


Figure is from d2l

## What we have done so far? Pretrained word embeddings

- The word embeddings can be pretrained by Skip-gram models or Glove models.

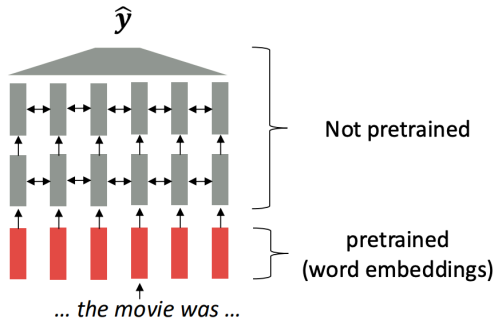
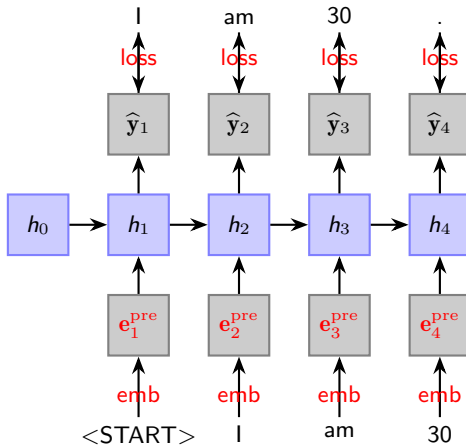


Figure is from Stanford CS 224n

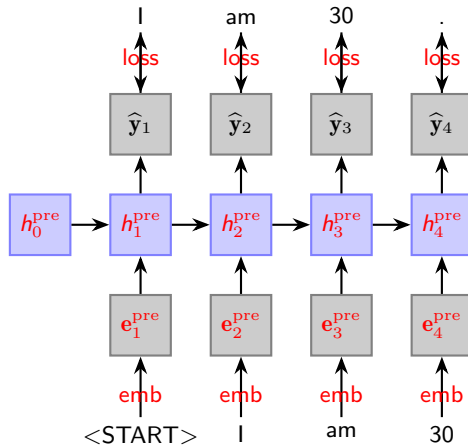
## What Have We Done So Far? Pretrained Word Embeddings (Assignment 3)

- Begin with pretrained word embeddings.
- Each word is assigned the same embedding, regardless of the sentence in which it appears.
- Use RNNs or Transformers to learn from context, updating embeddings based on the downstream task.
- The training data for downstream tasks (e.g., question answering, named entity recognition) must be sufficient to capture all relevant contextual information.

## Pretrained embeddings with glove



## Pretraining whole models are more popular in modern machine learning



## Pretraining and Fine-Tuning

- Let  $D_{\text{pretrain}}$  denote the dataset used for pretraining.
- Define the corresponding loss function  $\mathcal{L}_{\text{pretrain}}(\theta)$  based on  $D_{\text{pretrain}}$  and a deep neural network  $f_{\theta}$ .
- Learn the pretrained parameters  $\hat{\theta}_{\text{pretrain}}$  by minimizing the loss:

$$\min_{\theta} \mathcal{L}_{\text{pretrain}}(\theta).$$

- Obtain a foundation model  $f_{\hat{\theta}}$ .



## Pretraining and Fine-Tuning

- Let  $D_{\text{finetuning}}$  denote the downstream dataset, which can be relatively small.
- Define the corresponding loss function  $\mathcal{L}_{\text{finetuning}}(\theta)$  for a deep neural network  $f_{\theta}$ .
- Minimize  $\mathcal{L}_{\text{finetuning}}(\theta)$  using **SGD or its variants, initialized with the pretrained parameters  $\hat{\theta}_{\text{pretrain}}$** .

## Loss landscape of deep neural networks




Figure 1: The loss surfaces of ResNet-56 with/without skip connections. The proposed filter normalization scheme is used to enable comparisons of sharpness/flatness between the two figures.

Figure is from <https://arxiv.org/abs/1712.09913>

## Why Pretraining Matters: Possible Explanations

- Certain local minima near  $\hat{\theta}_{\text{pretrain}}$  may exhibit better generalization performance.
- The gradients of the fine-tuning loss near  $\hat{\theta}_{\text{pretrain}}$  may propagate more effectively, facilitating optimization.

## Sometimes We May Not Need to Fine-Tune All Parameters

- In some cases, we may **freeze** a subset of the pretrained parameters and only fine-tune the remaining ones.
- Mathematically, we partition the model parameters as  $\theta = (\theta^{(1)}, \theta^{(2)})$ .
- The pretrained parameters can be expressed as  $\hat{\theta}_{\text{pretrain}} = (\hat{\theta}_{\text{pretrain}}^{(1)}, \hat{\theta}_{\text{pretrain}}^{(2)})$ .
- To fine-tune only  $\theta^{(1)}$  while freezing  $\hat{\theta}_{\text{pretrain}}^{(2)}$ , we define the modified loss:

$$\mathcal{L}'_{\text{finetune}}(\theta^{(1)}) = \mathcal{L}_{\text{finetune}}(\theta^{(1)}, \hat{\theta}_{\text{pretrain}}^{(2)}).$$

- We then minimize this loss w.r.t.  $\theta^{(1)}$  using SGD or its variants, initializing at  $\hat{\theta}_{\text{pretrain}}^{(1)}$ .

## Linear Probing

- In fine-tuning, when only the last layer of a neural network is trainable while all preceding layers are frozen, the procedure is referred to as **linear probing** (since the model is linear).
- Linear probing is often sufficient when **the downstream dataset closely resembles a subset of the data used during pretraining**.
- This can be interpreted from the perspective of representation learning. Recommended reading:
  - Kumar et al. (2022). Fine-Tuning Can Distort Pretrained Features and Underperform Out-of-Distribution.
  - Wang et al. (2024). Neural Collapse Meets Differential Privacy: Curious Behaviors of NoisyGD with Near-Perfect Representation Learning.

# Linear Probing

- Recall the structure of a (fully connected) deep neural network:

$$f_{\theta}(\mathbf{x}) = \text{SoftMax} \left( \mathbf{W}^{[L]} \sigma(\mathbf{W}^{[L-1]} \sigma(\dots \sigma(\mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]})) + \mathbf{b}^{[L-1]}) + \mathbf{b}^{[L]} \right).$$

- Let  $\hat{\theta}_{\text{pretrain}} = (\widehat{\mathbf{W}}^{[L]}, \widehat{\mathbf{b}}^{[L]}, \dots, \widehat{\mathbf{W}}^{[1]}, \widehat{\mathbf{b}}^{[1]})$  denote the parameters obtained from pretraining.
- Define  $\widehat{\mathbf{h}}(\mathbf{x})$  as the output of the penultimate layer under the pretrained network, and consider the model:

$$f_{\mathbf{W}^{[L]}, \mathbf{b}^{[L]}}(\mathbf{x}) = \text{SoftMax} \left( \mathbf{W}^{[L]} \widehat{\mathbf{h}}(\mathbf{x}) + \mathbf{b}^{[L]} \right).$$

- In linear probing, only  $\mathbf{W}^{[L]}$  and  $\mathbf{b}^{[L]}$  are updated by minimizing the fine-tuning loss, while  $\widehat{\mathbf{h}}(\mathbf{x})$  remains fixed.

## Adding New Layers During Fine-tuning

- Recall that  $f_{\hat{\theta}}(\mathbf{x})$  denotes a pretrained neural network.
- To enhance the model's capacity or adapt it to a new task, we can append additional layers:

$$f_{\mathbf{W}^{\text{new}}, \mathbf{b}^{\text{new}}}(\mathbf{x}) = \mathbf{W}^{\text{new}} f_{\hat{\theta}}(\mathbf{x}) + \mathbf{b}^{\text{new}}.$$

- The loss function can then be minimized with respect to:
  - only the new parameters  $\mathbf{W}^{\text{new}}, \mathbf{b}^{\text{new}}$ , or
  - both the new parameters and the original pretrained parameters  $\theta$ .

## A Simple Fine-tuning Demo (Related to Final Project)

- A demonstration notebook is provided: `FineTuneBERT_mirror.ipynb`.



## Key Points to Keep in Mind When Fine-tuning

- Use the Hugging Face mirror. A tutorial is available [here](#).
- Use GPU acceleration—use our server for training.
- This is a warm-up for your final project:
  - Basic implementation similar to the demo (score range: 60–70).
  - Enhancements such as hyperparameter tuning, using larger or additional datasets, incorporating alternative models beyond BERT, or tackling more complex tasks with online datasets (score range: 70–80).
  - All the above, plus a clearly written and well-organized report (score range: 80–90).
  - All the above, with in-depth comparisons of models or datasets, the use of self-collected downstream datasets (including collection methodology), or the integration of advanced techniques (e.g., privacy-preserving methods, complex task settings such as multi-modal) (score range: 90–100).

① Section 1: An overview of pretraining and finetuning

② Section 2: Details of pretraining

③ Section 3: Details of GPT and BERT

## Key to Pretraining

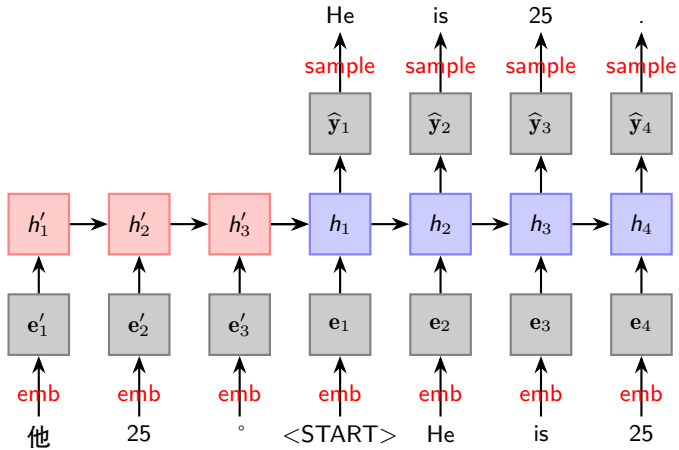
- The model should be sufficiently large to process large-scale, diverse datasets.
- Most of the data used in pretraining are unlabeled.
- The performance of pretrained models often follows a *scaling law*, where increased data, model size, and compute lead to improved results.

## Why Unlabeled Data?

- The vast majority of data available is unlabeled.

Dataset	Number of Tokens
SQuAD 2.0 (labeled)	< 50 million
DCLM-pool	240 trillion
Pretraining Data	510T (indexed), 3100T (total)

## Recap: The Encoder–Decoder Architecture



## Three Approaches to Pretraining


- **Encoder-only:** Trained with bidirectional context—can "see" both past and future tokens (e.g., BERT).
- **Decoder-only:** Autoregressive language models that predict the next token based only on past context (e.g., GPT).
- **Encoder-decoder:** Combines both components—encoder processes the input, decoder generates the output (e.g., T5).

## Pretraining Encoders

- Encoders are trained to capture bidirectional context—but how?
- Mask a subset of tokens using a special [MASK] token during training.
- For example: “I [MASK] to a store.”
- This model is called a **masked language model**

# Pretraining Encoders

$\ell(\theta) = -\log$  predicted probability of “went”





## Pretraining Encoders: Masked Language Modeling

- Given a sequence  $\mathbf{x} = (x_1, \dots, x_L)$ , let  $h_t = \text{Encoder}(x_t)$  denote the contextual embedding at position  $t$ .
- Randomly replace a subset of tokens  $\{x_{t_i}\}_{i=1}^{L'}$  with the special token [MASK].
- Predict the original tokens using:

$$\hat{y}_{t_i} = \text{SoftMax}(\mathbf{W}_y h_{t_i}([\text{MASK}]) + \mathbf{b}_y).$$

- Compute the cross-entropy loss using the ground truth label, e.g.,  $y_2 = \text{"went"}$ :

$$\text{CE}(\hat{y}_{t_i}, y_{t_i}).$$

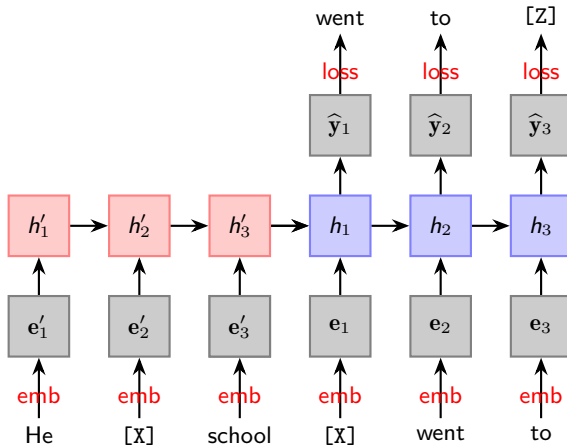
- The final loss is the average over all masked positions:

$$\frac{1}{L'} \sum_{i=1}^{L'} \text{CE}(\hat{y}_{t_i}, y_{t_i}).$$

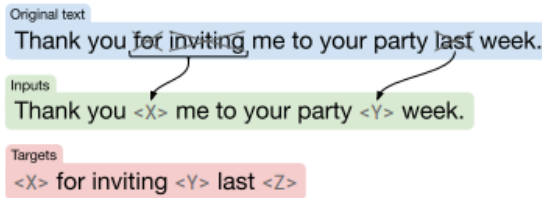
## Pretraining Encoder-Decoder Models

- We illustrate the use of the “Text-to-Text Transfer Transformer” (T5), a pretrained encoder-decoder model proposed by Raffel et al., 2018.
- **Encoder training:** The input sequence is partially **masked** by replacing **variable-length spans** with unique sentinel tokens.
- **Decoder training:** The target output is formed by concatenating the dropped spans in order, each delimited by the corresponding sentinel token, ending with a final sentinel token [Z].

## Pretraining encoder-decoder models



## Pretraining encoder-decoder models




## Pretraining encoder-decoder models

- The encoder-decoder model may have better performance in some tasks like denoising.

Architecture	Objective	Params	Cost	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Encoder-decoder	Denoising	$2P$	$M$	<b>83.28</b>	<b>19.24</b>	<b>80.88</b>	<b>71.36</b>	<b>26.98</b>	<b>39.82</b>	<b>27.65</b>
Enc-dec, shared	Denoising	$P$	$M$	82.81	18.78	<b>80.63</b>	<b>70.73</b>	26.72	39.03	<b>27.46</b>
Enc-dec, 6 layers	Denoising	$P$	$M/2$	80.88	18.97	77.59	68.42	26.38	38.40	26.95
Language model	Denoising	$P$	$M$	74.70	17.93	61.14	55.02	25.09	35.28	25.86
Prefix LM	Denoising	$P$	$M$	81.82	18.61	78.94	68.11	26.43	37.98	27.39
Encoder-decoder	LM	$2P$	$M$	79.56	18.59	76.02	64.29	26.27	39.17	26.86
Enc-dec, shared	LM	$P$	$M$	79.60	18.13	76.35	63.50	26.62	39.17	27.05
Enc-dec, 6 layers	LM	$P$	$M/2$	78.67	18.26	75.32	64.06	26.13	38.42	26.89
Language model	LM	$P$	$M$	73.78	17.54	53.81	56.51	25.23	34.31	25.38
Prefix LM	LM	$P$	$M$	79.68	17.84	76.87	64.86	26.28	37.51	26.76

## Pretraining decoders



## Pretraining decoders

- Given a sequence  $\mathbf{x} = (x_1, \dots, x_L)$ , let  $h_t = \text{Encoder}(x_t)$  denote the contextual embedding at position  $t$ .
- Predict the  $t + 1$ -th token with probability:

$$\hat{y}_t = \text{SoftMax}(\mathbf{W}_y h_t(x_t) + \mathbf{b}_y).$$

- Compute the cross-entropy loss using the ground truth label

$$\text{CE}(\hat{y}_t, y_t).$$


- The final loss is the average over all masked positions:

$$\frac{1}{L} \sum_{t=1}^L \text{CE}(\hat{y}_t, y_t).$$





# Recap:transformer encoder decoder



# Bidirectional Encoder Representations from Transformers (BERT)

## **BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding**

**Jacob Devlin   Ming-Wei Chang   Kenton Lee   Kristina Toutanova**

**Google AI Language**

`{jacobdevlin, mingweichang, kentonl, kristout}@google.com`

- Each input sentence in BERT's training data is wrapped with two special tokens: [CLS] at the beginning and [SEP] at the end (or between two segments).
- For example:  

`[CLS] my dog is cute [SEP] he likes play ##ing [SEP]`
- The [CLS] token is used for classification tasks, while [SEP] marks the end of a sentence or separates two segments.

## Input Embeddings

- Token Embeddings: Pretrained WordPiece embeddings (Wu et al., 2016) with a vocabulary of 30,000 tokens.
- Positional Embeddings: As introduced in the Transformer section, these embeddings encode the position of tokens in a sequence.
- Segment Embeddings: These embeddings differentiate between sentences.
- Input Embeddings:  $e_t = \text{Token Embeddings} + \text{Positional Embeddings} + \text{Segment Embeddings}$

# Input Embeddings




Figure is from Devlin et al., 2018

## Masked Language Modeling for BERT

- In this pretraining task, **15% of tokens** are randomly selected to be masked for prediction.
- However, the special token [Mask] is never used during fine-tuning.
- To prevent a mismatch between pretraining and fine-tuning, if a token is masked for prediction, it is replaced with other tokens **randomly** as follows:
  - A special “<mask>” token is used 80% of the time (e.g., “this movie is great” becomes “this movie is <mask>”).
  - The original token is used 10% of the time (e.g., “this movie is great” remains “this movie is great”).
  - A random token is used 10% of the time (e.g., “this movie is great” becomes “this movie is drink”).

## Next Sentence Prediction

- While masked language modeling captures bidirectional context to represent words, it does not explicitly model the logical relationship between text pairs.
- To address this, BERT incorporates a binary classification task called next sentence prediction during its pretraining.
- During pretraining, half of the sentence pairs are consecutive sentences labeled as “True.”
- For the other half, the second sentence is randomly sampled from the corpus and labeled as “False.”
- This task is facilitated by the [CLS] token. What is the embedding of the [CLS] token?

## Next Sentence Prediction

- The [CLS] token has only two possible values (True or False), and its embedding  $e_{[\text{CLS}]}$  lies in  $\mathbb{R}^2$ .

```
class NextSentencePred(nn.Module):
    """The next sentence prediction task of BERT."""
    def __init__(self, **kwargs):
        super(NextSentencePred, self).__init__(**kwargs)
        self.output = nn.Linear(2)

    def forward(self, X):
        # `X` shape: (batch size, `num_hiddens`)
        return self.output(X)
```



## Model Structure and Details of BERT-Base

- We begin by introducing a smaller model—BERT-base.
- Each encoder transformer block (ignoring the "Add & Norm" layer) has the following structure:  
FF(Multi-head attention(Input))
- Output dimension of the multi-head attention layer: 768-dimensional hidden states.
- Output dimension of the feed-forward layer:  $4 \times 768$ .
- Number of attention heads: 12.
- Depth: 12 encoder transformer blocks.

## Model Structure and Details of BERT-Large

- BERT-Large is a larger model than BERT-base.
- Output dimension of the multi-head attention layer: 1024-dimensional hidden states.
- Output dimension of the feed-forward layer:  $4 \times 1024$ .
- Number of attention heads: 16 (The dimension of each head is the same as BERT-base).
- Depth: 16 encoder transformer blocks.

## Training data for BERT

- BooksCorpus (800 million words)
- English Wikipedia (2,500 million words)

## Pretraining is time-consuming

- BERT is pretrained on TPUs, specialized hardware designed to accelerate tensor operations.
- The pretraining of BERT was conducted using 64 TPU chips over a span of 4 days.
- Finetuning is practical and common on a single GPU!!!

## Fine-Tuned BERT is Powerful on Many Diverse Datasets

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT <sub>BASE</sub>	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT <sub>LARGE</sub>	<b>86.7/85.9</b>	<b>72.1</b>	<b>92.7</b>	<b>94.9</b>	<b>60.5</b>	<b>86.5</b>	<b>89.3</b>	<b>70.1</b>	<b>82.1</b>

Figure adapted from Devlin et al., 2018

# The Larger, the Better

System	Dev		Test	
	EM	F1	EM	F1
Top Leaderboard Systems (Dec 10th, 2018)				
Human	-	-	82.3	91.2
#1 Ensemble - nlnet	-	-	86.0	91.7
#2 Ensemble - QANet	-	-	84.5	90.5
Published				
BiDAF+ELMo (Single)	-	85.6	-	85.8
R.M. Reader (Ensemble)	81.2	87.9	82.3	88.5
Ours				
BERT <sub>BASE</sub> (Single)	80.8	88.5	-	-
BERT <sub>LARGE</sub> (Single)	84.1	90.9	-	-
BERT <sub>LARGE</sub> (Ensemble)	85.8	91.8	-	-
BERT <sub>LARGE</sub> (Sgl.+TriviaQA)	<b>84.2</b>	<b>91.1</b>	<b>85.1</b>	<b>91.8</b>
BERT <sub>LARGE</sub> (Ens.+TriviaQA)	<b>86.2</b>	<b>92.2</b>	<b>87.4</b>	<b>93.2</b>

Table 2: SQuAD 1.1 results. The BERT ensemble is 7x systems which use different pre-training checkpoints and fine-tuning seeds.

Figure adapted from Devlin et al., 2018

## Limitations of BERT

- So, why don' t we use BERT for everything?
- BERT, like other pretrained encoders, is not naturally suited for autoregressive sequence generation tasks.
- If your task involves sequence generation, consider using a pretrained decoder.

## Generative Pretrained Transformer (GPT)

- GPT-2 (Radford et al., 2018) was a significant success in pretraining a decoder!
- Transformer decoder with 12 layers and 117M parameters.
- 768-dimensional hidden states and 3072-dimensional feed-forward hidden layers.
- Tokens and Embeddings: Pretrained Byte-Pair Encoding (BPE) with 40,000 merges (tokens).
- Trained on BooksCorpus: over 7,000 unique books, containing long spans of contiguous text to learn long-distance dependencies.